

CSCI 3753

Operating Systems

Virtual Machines

Chapters 16

Lecture Notes By

Shivakant Mishra

Computer Science, CU-Boulder

Last Update: 12/07/17

Terminology

- “Virtual machine” is a loaded term
 - E.g. Java Virtual Machine refers to a runtime environment (software) that can execute Java bytecode
- “VM” is a loaded abbreviation
 - JVM (Java Virtual Machine), Virtual Memory
- For our purposes, we will talk about Virtual Machine Monitors (VMM)
 - VMM is software that allows multiple guest OSes to run concurrently on one physical machine
 - Each guest runs on a virtual machine
 - VMM is sometimes called a *hypervisor*

Virtual Machine Monitors (VMMs)

- VMMs are everywhere
- Industry commitment
 - Software: Vmware, Xen, ...
 - Hardware: Intel-VT, AMD-V
 - If Intel and AMD add it to their chip, you know it's serious ...
- An old idea, actually developed by IBM in 60's

What is a VMM?

- We have seen that an OS already virtualizes
 - Syscall, processes, virtual memory, file system, sockets, ...
 - Applications program to this interface
- A process already is given the illusion that it has its
 - Own memory, via virtual memory
 - Own CPU, via time slicing
 - Own I/O devices, via device-independent I/O

What is a VMM?

- A VMM extends this idea and virtualizes an entire physical machine
 - Interface supported is the hardware
 - OS defines a higher level interface
 - VMM provides the illusion that software has full control over the hardware (of course, VMM is in control)
 - VMM “applications” run in virtual machines
- Implications
 - You can boot an OS in a VM
 - Run multiple instances of an OS on same physical machine
 - Run different OSes simultaneously on the same machine

Why VMMs?

- Resource utilization
 - Machines today are powerful, want to multiplex their hardware
 - E.g. Cloud providers
 - Can migrate VMs from one machine to another without shutdown
- Software use and development
 - Can run multiple OSes simultaneously
 - E.g. No need to dual boot
 - Can do system development at the user level
- Many other cool applications
 - Debugging, emulation, security, speculation, fault tolerance
- Common theme is manipulating applications/services at the granularity of a machine

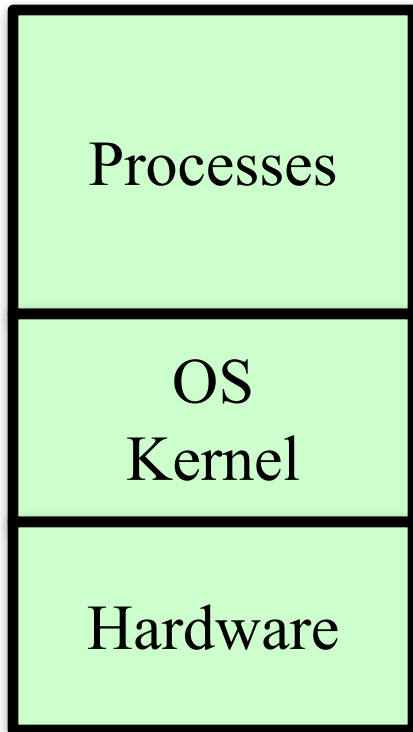
Goals of Virtualization

- Fidelity
 - OSes and applications work the same without modification
 - Except timings ...
 - Although we may modify the OS a bit
- Isolation
 - VMM protects resources and VMs from each other
- Performance
 - VMM is another software layer --- overhead
 - VMware
 - CPU-intensive apps: 2-10% overhead
 - I/O intensive apps: 25-60% overhead
 - An overwhelming majority of guest instructions are executed by the hardware without VMM intervention

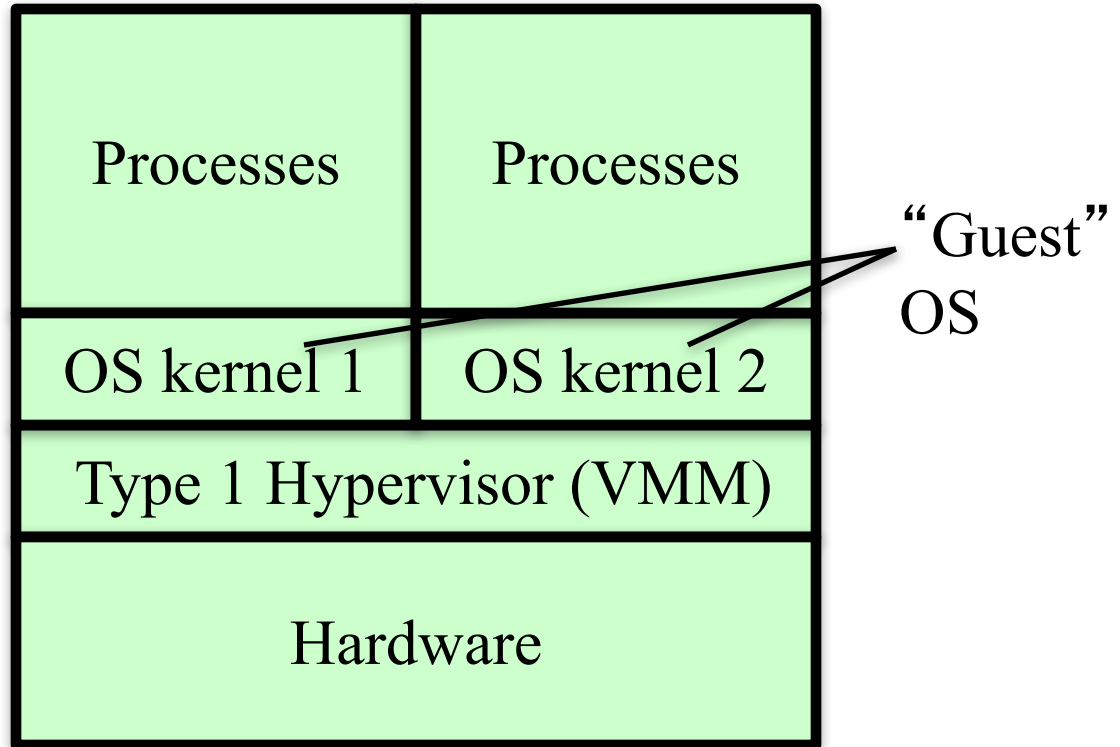
Announcements

- Final Exam
 - Monday, 12/18, 4:30-6:00 PM in class
 - Special accommodation: Monday, 12/18, 4:30 - ... in ECES 112 C.
 - Material
 - Memory management, mass storage structure, file systems, security and protection, virtual machines
 - Chapters 8, 9, 10, 11, 12, 14, 15 and 16
 - Lecture sets 12, 13, 14, 15, 16, 17, 18, 19 and 20
 - One question from midterm exam
- Format
 - Similar to midterm exam
 - Closed notes, closed book, ...

Type 1 Hypervisor



Traditional OS



Type 1 Hypervisor: runs directly on HW
Example: IBM's CP/CMS, Citrix XenServer,
VMware ESXi, Microsoft Hyper-V

Type 1 Hypervisor

- *Type 1* (or *native*) hypervisors run directly on the host's hardware
 - A guest operating system thus runs on another level above the hypervisor
- Also known as Full virtualization
- The hypervisor provides illusion of identical hardware as the real system
 - Guest OS doesn't have to be recompiled, and doesn't even realize it's executing on a VM
 - Example VMs: Microsoft Virtual Server, VMWare ESX Server
 - Disadvantage: Virtualized applications can run much slower, because of the many layers of fully virtualized abstraction between them and the hardware

Full Virtualization


- There are two approaches in full virtualization
 - Hardware assisted virtualization
 - Provides architectural support that facilitates building a virtual machine monitor and allows guest OSs to be run in isolation
 - Software assisted virtualization
 - the virtual machine simulates enough hardware to allow an unmodified "guest" OS to be run in isolation.

Type 1 Hypervisor: How it Works

- VMM can interpret every instruction from guest OSes and execute them
- Problem: emulating/interpreting every instruction are not good options – too much software overhead
- Goal: Want to create a virtual machine that executes at close to native speeds on a CPU
- Solution: Have the guest OS execute normally, directly on the CPU, except that it is not in kernel mode.
 - Any privileged instructions invoked by the guest OS will be trapped to the hypervisor, which is in kernel mode.

Protected Mode

- Most modern CPUs support protected mode
- x86 CPUs support three rings with different privileges
 - Ring 0: OS kernel
 - Ring 1, 2: device drivers
 - Ring 3: userland
- Most OSes only use rings 0 and 3



Challenges With Virtual Hardware

Issue: x86 is not designed with virtualization in mind

1. Dealing with privileged instructions

- OSes expect to run with high privilege (ring 0)
- How can the VMM enable guest OSes to run in user mode (ring 3)?

2. Managing virtual memory

- OSes expect to manage their own page tables
- MMU supports only one level of virtualization
- How can the VMM translate between a guest's page tables and the hosts page tables?

Privileged Instructions

- OSes rely on many privileges of ring 0
 - `cli`, `sti`, `popf` – Enable/disable interrupts
 - `hlt` – Halt the CPU until the next interrupt
 - `mov cr3, 0x00FA546C` – install a page table
 - Install interrupt and trap handlers
 - Etc...
- However, guest OSes run in user mode
- VMM must somehow virtualize privileged operations

Using Exceptions for Virtualization

- Ideally, when a guest executes a privileged instruction in ring 3, the CPU should generate an exception
- Example: suppose the guest executes `hlt`
 1. The CPU generates a protection exception
 2. The exception gets passed to the VMM
 3. The VMM can emulate the privileged instruction
 - If guest 1 runs `hlt`, then it wants to go to sleep
 - VMM can do `guest1.yield()`, then schedule guest 2

Problem: x86 Doesn't Except Properly

Binary Translation

- x86 assembly cannot be virtualized because some privileged instructions don't generate exceptions
- Workaround: translate the unsafe assembly from the guest to safe assembly
 - Known as binary translation
 - Performed by the VMM
 - Privileged instructions are changed to function calls to code in VMM

Type 1 Hypervisor: How it Works

- The hypervisor then emulates only these privileged instructions and when done passes control back to the guest OS, also known as a “VM entry”
- This way, most ordinary (non-privileged) instructions operate at full speed, and only privileged instructions incur the overhead of a trap, also known as a “VM exit”, to the hypervisor/VMM.
- This approach to VMs is called *trap-and-emulate*

Binary Translation Example

Guest OS Assembly

Translated Assembly

do_atomic_operation:

cli

mov eax, 1

xchg eax, [lock_addr]

test eax, eax

jnz spinlock

...

...

mov [lock_addr], 0

sti

ret

do_atomic_operation:

call [vmm_disable_interrupts]

mov eax, 1

xchg eax, [lock_addr]

test eax, eax

jnz spinlock

...

...

mov [lock_addr], 0

call [vmm_enable_interrupts]

ret

Binary Translation: Pros and Cons

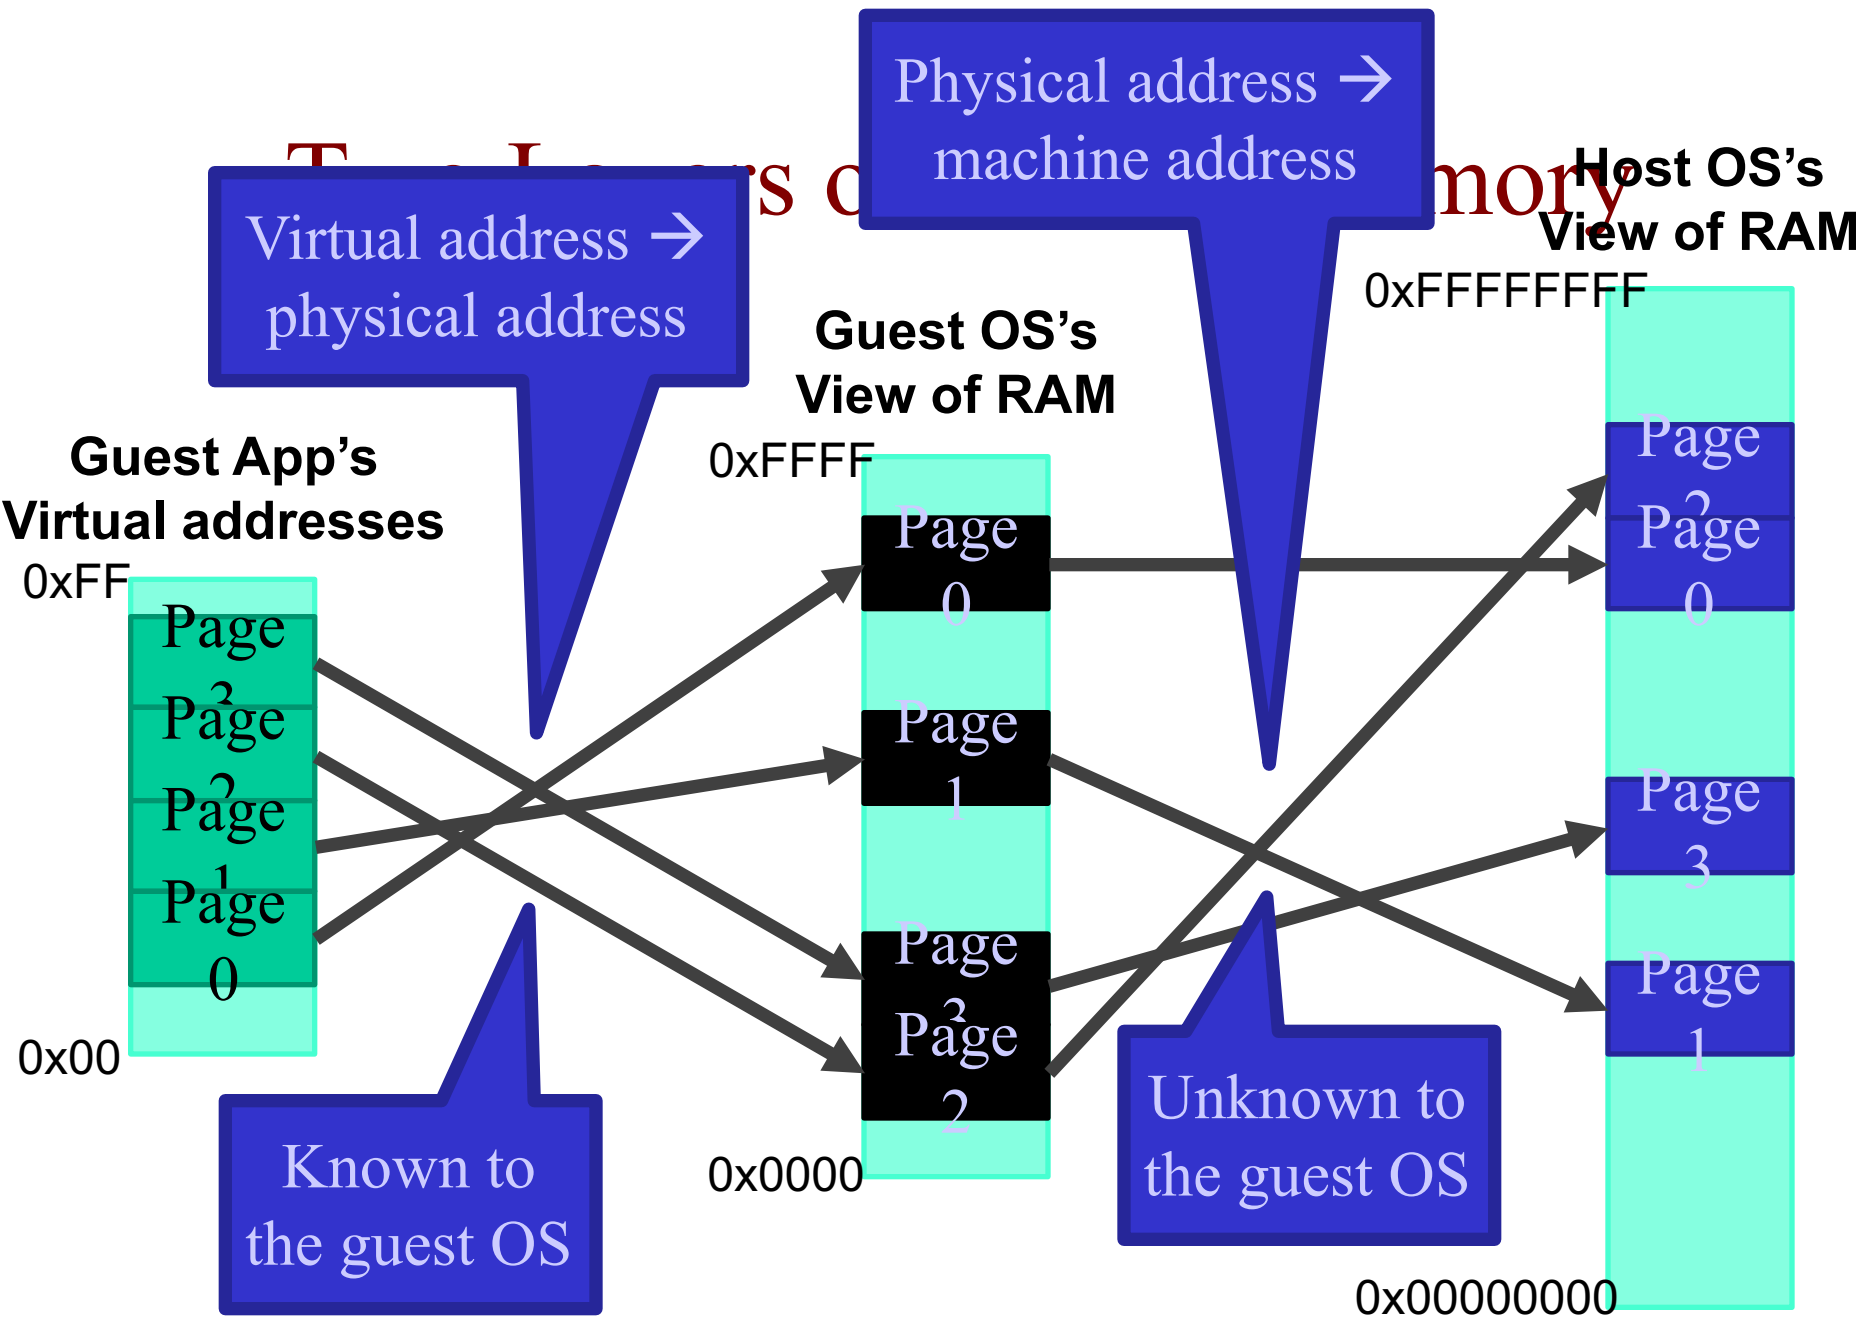
- Advantages of binary translation
 - It makes it safe to virtualize x86 assembly code
 - Translation occurs dynamically, on demand
 - No need to translate the entire guest OS
 - App code running in the guest does not need to be translated
- Disadvantages
 - Translation is slow
 - Wastes memory (duplicate copies of code in memory)
 - Translation may cause code to be expanded or shortened
 - Thus, jmp and call addresses may also need to be patched

Caching Translated Code

- Typically, VMMs maintain a cache of translated code blocks
 - LRU replacement
- Thus, frequently used code will only be translated once
 - The first execution of this code will be slow
 - Other invocations occur at native speed

Problem: How to Virtualize the MMU?

- On x86, OS expects that it can create page tables and install them in the *cr3* register
 - The OS believes that it can access physical memory
- However, virtualized guests do not have access to physical memory
- Using binary translation, the VMM can replace writes to *cr3*
 - Store the guest's root page in the virtual CPU *cr3*
 - The VMM can now walk to guest's page tables
- However, the guest's page tables cannot be installed in the physical CPU...

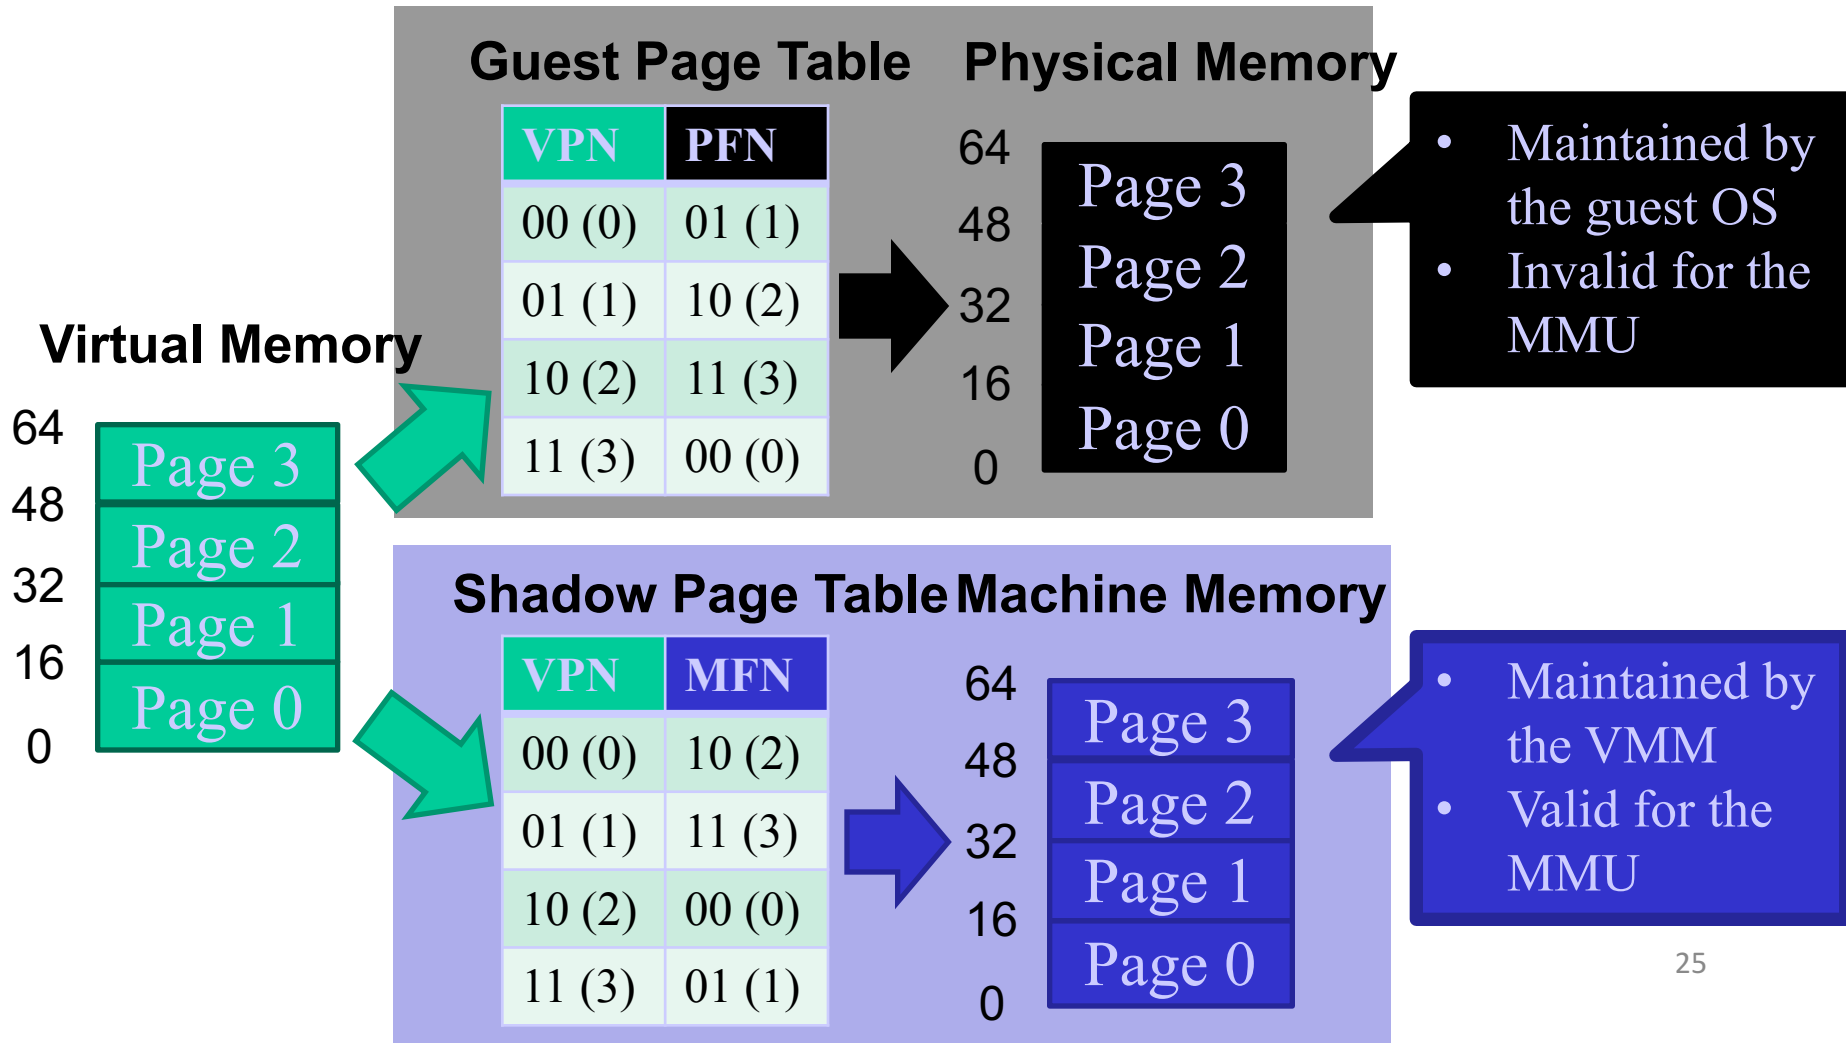


Guest's Page Tables Are Invalid

- Guest OS page tables map virtual page numbers (VPNs) to physical frame numbers (PFNs)
- Problem: the guest is virtualized, doesn't actually know the true PFNs
 - The true location is the machine frame number (MFN)
 - MFNs are known to the VMM
- Guest page tables cannot be installed in *cr3*
 - Map VPNs to PFNs, but the PFNs are incorrect
- How can the MMU translate addresses used by the guest (VPNs) to MFNs?

Shadow Page Tables

- Solution: VMM creates shadow page tables that map VPN → MFN (as opposed to VPN → PFN)



Building Shadow Tables

- Problem: how can the VMM maintain consistent shadow pages tables?
 - The guest OS may modify its page tables at any time
 - Modifying the tables is a simple memory write, not a privileged instruction
 - Thus, no helpful CPU exceptions :(
- Solution: mark the hardware pages containing the guest's tables as read-only
 - If the guest updates a table, an exception is generated
 - VMM catches the exception, examines the faulting write, updates the shadow table
- Special attention needs to be paid to page fault handling

Shadow page tables: Pros and Cons

- The good: shadow tables allow the MMU to directly translate guest VPNs to hardware pages
 - Thus, guest OS code and guest apps can execute directly on the CPU
- The bad:
 - Double the amount of memory used for page tables
 - i.e. the guest's tables and the shadow tables
 - Context switch from the guest to the VMM every time a page table is created or updated
 - Very high CPU overhead for memory intensive workloads

Hardware Techniques

- Modern x86 chips support hardware extensions designed to improve virtualization performance
 1. Reliable exceptions during privileged instructions
 - Known as AMD-V and VT-x (Intel)
 - Released in 2006
 - Adds vmrun/vmexit instructions (like sysenter/sysret)
 2. Extended page tables for guests (second level addr translation)
 - Known as RVI (AMD) and EPT (Intel)
 - Released in 2008
 - Adds another layer onto existing page table to map PFN→MFN

Pros and Cons of AMD-V and VT-x

- Greatly simplifies VMM implementation
 - No need for binary translation
 - Simplifies implementation of shadow page tables
- ... however, sophisticated VMMs still use binary translation in addition to `vmenter/vmexit`
 - Some operations are **much** slower when using *vmexit* vs. binary translation
 - VMM observes guest code that causes frequent `vmexits`
 - Hot spots may be binary translated or dynamically patched to improve performance

Pros and Cons of RVI and EPT

- Huge performance advantages vs. shadow page tables
- Memory overhead
 - ... but not as much as shadow page tables
- TLB misses are twice as costly
 - page tables twice as deep, hence it takes twice as long to resolve TLB misses

Configuring Your VMM

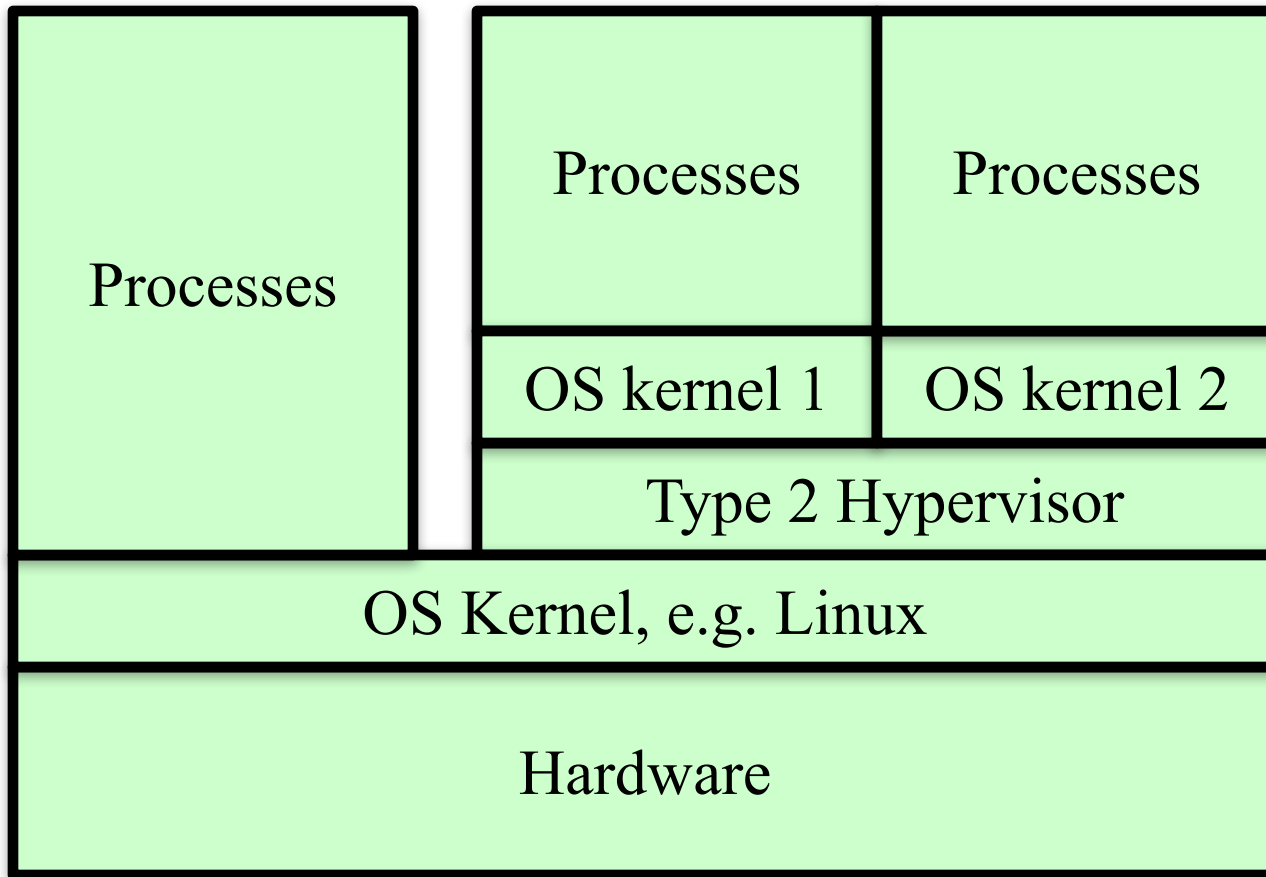
- Advanced VMMs like VMWare give you three options
 1. Binary translation + shadow page tables
 2. AMD-V/VT-x + shadow page tables
 3. AMD-V/VT-x + RVI/EPT
- Which is best?
 - Choosing between 1 and 2 is more difficult
 - For some workloads, 2 is much slower than 1
 - Run benchmarks with your workload before deciding on 1 or 2

- Fastest by far
- But, requires very recent, expensive CPUs

Type 2 Hypervisor

- Problem: (Full virtualization) it takes a lot of work to virtualize an arbitrary guest OS
 - VMM implementation is very complicated
 - Even with hardware support, performance issues remain
- What if we require that guests be modified to run in the VMM
 - How much work is it to modify guests to “cooperate” with the VMM?
 - Will VMM implementation be simpler?
 - Can we get improved performance?

Type 2 Hypervisor



Type 2 Hypervisor: runs on an OS
Example: KVM, Virtualbox

Type 2 Hypervisor

- *Type 2 (or hosted)* hypervisors run within a conventional operating system environment
 - The hypervisor layer as a distinct 2nd software level
 - A guest operating system runs at the third level
- There are two approaches
 - Operating system-level virtualization
 - Allows multiple isolated and secure virtualized servers to run on a single physical server.
 - The OS kernel is used to implement the "guest" environment
 - Applications running in a given "guest" environment view it as a stand-alone system
 - Paravirtualization

Paravirtualization

- *Paravirtualization* addresses the performance issue:
 - Paravirtualization provides specially defined 'hooks' to allow the guest(s) and host to request and acknowledge certain tasks, which would otherwise be executed in the hypervisor
 - An OS no longer sees full virtualization of hardware, but instead must be specially recompiled to work with this particular hypervisor, exploiting special optimizations to speed performance

Paravirtualization

- Involves modifying the OS to run in the virtualized environment as a VM → needs to be recompiled
- Virtual machine does not necessarily simulate hardware, but instead (or in addition) offers a special API that can only be used by modifying the "guest" OS
 - VMware tries to run code whenever possible directly on CPU without emulation, e.g. user code
- Example VMs: Xen, Parallels, VMware Workstation

Hypercalls

- The Xen VMM exports a hypercall API
 - Methods replace privileged instructions offered by the hardware
 - E.g halt CPU, enable/disable interrupts, install page table
 - Guest OS can detect if it's running directly on hardware or on Xen
 - In the former case, typical ring 0 behavior is used
 - In the latter case, hypercalls are used
- If a guest executes a privileged instruction, crash it
 - Xen VMM makes no attempt to emulate privileged instructions
 - Simplifies Xen VMM implementation

Modifying Guests

- How much work does it take to modify a guest OS to run on Xen?
 - Linux: 3000 lines (1.36% of the kernel)
 - Including device drivers
 - Windows XP: 4620 lines (0.04% of the kernel)
 - Device drivers add another few hundred lines of code
- Modification isn't trivial, but its certainly doable

Virtualization

- Virtualization has made a huge resurgence in the last 20 years
- Today, all OSes and most CPUs have direct support for hosting virtual machines, or becoming virtualized
- Virtualization underpins the cloud
 - E.g. Amazon EC2 rents virtual machines at low costs
 - Hugely important for innovation